



Technical University of Denmark



High Performance Computing

FORTRAN, OpenMP and MPI

41391

Content

- Day 5:
 - Types of files in HPC.
 - Data transfer.
 - Operations on external files.
 - Recommendations.
 - Submitting jobs (batch systems).
 - Memory and caches.
 - Revision control systems (git)
 - Fortran 2003/2008 extensions.

Types of files in HPC

- Setup file: small ASCII file.
- Input files: large binary/ASCII files.
- Post processing files: single precision snapshots in time/iterations (FORTRAN binary, NetCDF, HDF5, VTK, PLOT3D, ...)
- Diagnostic files (small ASCII files; maybe binary + small utility program to convert to ASCII).
- Restart files: full precision image of the memory/state of the problem.
- Auxiliary files.

Data transfer

Conversion

- Data transfer to external files (data outside the main memory) requires a conversion of the internal machine representation (e.g., hexadecimal) to a human readable (ASCII) format.
- The conversion is performed according to an *edit descriptor* contained in a *format description*.

Example:

Convert (from hexadecimal) and write the number with the value 0.0045 use the edit descriptor: F10.4

(10 is the total width of the field, and 4 the length of the fractional part).

Conversion

- The format description is a list of edit descriptors enclosed in parentheses.

Example:

```
WRITE(*,'(I10,F10.3,A10)') ...
```

The format description is a character string



or in a separate format statement:

```
WRITE(*,10) ...
```

```
10 FORMAT(I10,F10.3,A10)
```

- Use READ and PRINT/WRITE for input and output.

PRINT only produces output to the screen;
WRITE also allows IO to external files

I/O lists and format definition

- The quantities to be read or written by a program are specified in an *I/O list*. For output they may be expressions, but for input they must be variables.

Example: to print the integer i , the real f , and the character c use the PRINT statement:

```
PRINT '(I10,F10.3,A10)',i,f,c
```

```
PRINT 10,i,f,c
```

```
10  FORMAT(I10,F10.3,A10)
```

Advice: use PRINT for easy/short debugging print; and/but use WRITE for permanent IO

I/O lists and format definition

Example: INPUT

INTEGER :: j(3)

CHARACTER(LEN=256) :: form = '(3I10)'

READ '(3I10)', j

READ form, j

READ 10, j

READ*,j! *list-directed I/O*

10 FORMAT(3I10)

Advice: always use list-directed I/O (the '*' syntax) for ASCII files and terminal IO.

I/O lists and format definition

Example: OUTPUT

```
INTEGER :: i,ilen
```

```
REAL, DIMENSION(10) :: a
```

```
CHARACTER(LEN=20) :: word
```

```
PRINT '(I10)',i
```

```
PRINT '(F10.3)', 3.0+5.0+SQRT(34.)
```

```
PRINT '(10F10.3)', a
```

```
PRINT '(3F10.3)', a(1:3) ! or (a(i), i=1,3)
```

```
PRINT '(A,A10)', 'Here is a word: ',word
```

```
ilen = LEN_TRIM(word)
```

```
PRINT '(2A)', 'Here is a word: ',word(1:ilen)
```

Unit numbers

- I/O is associated with a UNIT number, a default scalar integer in the range 1 – 99:

- READ(*u*,*fmt*) *list*
- WRITE(*u*,*fmt*) *list*

STDERR: unit = 0

STDIN: unit = 5

STDOUT: unit = 6

Example:

! Unit number:

READ(4,*) q

! Scalar default integer *constant*

READ(nunit,*) q

! Scalar default integer *variable*

READ(i*4+j,*) q

! Scalar default integer *expression*

READ(*,*) q

! The first asterisk (*) implies IO to
! the STDIN/STDOUT

Internal files

- *Internal files* allow format conversion between various representations to be carried out by the program in a storage area defined within the program itself.

Example:

```
CHARACTER(LEN=256) :: filename
```

```
INTEGER :: istep
```

```
WRITE(filename,'(A,I6.6,A)') 'post',istep,'.dat'
```

```
OPEN(10,FILE=filename)
```

```
WRITE(10,*) ...
```

I6.6: set aside 6 digits for the number with leading zeros: fx.: post000100.dat

Internal files

Example (*scanning* a file):

```
CHARACTER(LEN=256) :: buffer
```

```
INTEGER :: freq
```

```
READ(10,'(A)') buffer
```

```
IF (buffer(1:1).NE.'#') THEN ! Skip comments
```

```
    READ(buffer,*) freq ! Read/convert from character to integer
```

```
ENDIF
```

Formatted input

- The READ statement:

```
READ([UNIT=]u,[FMT=]fmt[,IOSTAT=ios] &  
[,ERR=error-label][,END=end-label]) list
```

Example:

```
i = 1
```

```
DO
```

```
    READ(10,*,IOSTAT=ios,ERR=100,END=200) buffer(i)
```

```
    i = i + 1
```

```
ENDDO
```

```
100 CONTINUE ! Error label
```

```
    WRITE(*, '(2(A,I5))') 'Error reading file in line: ', i, ' IOSTAT = ', ios
```

```
200 CONTINUE ! End of file label
```

```
    WRITE(*, '(A,I6,A)') 'Read ', i-1, ' lines'
```

Formatted output

- The WRITE statement:

WRITE([UNIT=]u,[FMT=]fmt[,IOSTAT=ios] &
[,ERR=error-label][,END=end-label]) list

- The components are similar to the READ.
- If the WRITE uses a unit number that is not opened (using the OPEN(...)) a default file name will be used (fx. ftn<u> or fort.<u>).

List-directed I/O

- List-directed I/O (asterisk format) will use a format defined by the computer system.
- WRITE of long lines may break the line. Hence a corresponding READ of the same line (with a different compiler) might fail!

List-directed I/O

Example:

```
READ(*,*) cbuf1
```

```
READ(*,'(A)') cbuf2
```

Will read the two lines:

```
"FORTRAN 90"
```

```
FORTRAN 90
```

As:

```
cbuf1 = FORTRAN 90
```

```
cbuf1 = FORTRAN
```

```
cbuf2 = "FORTRAN 90"
```

```
cbuf2 = FORTRAN 90
```


List-directed I/O

- Separators: blank, comma or slash (/).
- If the length of the variable is longer than the string, blanks are padded at the end.

Example:

```
CHARACTER(LEN=256) :: cbuf = 'test'
```

```
i = LEN_TRIM(cbuf)
```

```
WRITE(*,'(A)') cbuf ! will write 256 characters
```

```
WRITE(*,'(A)') cbuf(1:i) ! will write 4 characters
```

Namelist I/O

- *Namelist* is an annotated I/O list.

Example:

```
INTEGER :: Nx = 20, Ny = 30
```

```
REAL :: diff, array(5)
```

```
CHARACTER(LEN=20) :: name = 'string'
```

```
NAMELIST /list/ Nx, Ny, diff, array, name
```

```
WRITE(*, nml=list)
```

```
READ(*, nml=list)
```

Namelist I/O

will write/read the file:

&LIST

NX = 20

NY = 30

DIFF = 0.0

ARRAY = 0.0E+0 0.0E+0 0.0E+0 0.0E+0 0.0E+0

NAME = string

/

- Any element can be left out !
- In any multiple occurrence of an object for namelist input, the final value is taken.
- Lines beginning with ! are comment lines.

Non-advancing I/O

- Non-advancing input and output through the ADVANCE='no' (or 'yes')

Example:

```
WRITE(*,'(A)',ADVANCE='NO') 'Enter dt '
```

```
READ(*,*) dt
```

Edit descriptors

- Three classes:
 - Data (formatting/conversion).
 - Control (control of leading sign, blank separators).
 - Character-string.
- The data edit descriptors may have repeat counters.

Examples:

```
WRITE(*,'(4(I5,F8.2))') (i(j),a(j),j=1,4)
```

Edit descriptors

- **I**NTEGER: i-descriptor:
 - iw.m: w = width of the field and m is the least number of digits to be output.

Example:

The number: 99 printed with: i5.3 returns: bb099
- For the *i* and all the other numerical edit descriptors, if the output field is too narrow to contain the number to output, it is filled with asterisks (*****).

Edit descriptors

- REAL: e,en,es or f-descriptors:
 - fw.d: w = width of the field and d is the number of decimal points after the comma.
 - ew.d and ew.dee (latter for large exponents).
 - enw.d: (engineering) As the e -format, but the decimal exponent is divisible with three (3).
 - esw.d: (scientific): as the e -format, but the nonzero significant is greater or equal to one and less than 10.
 - Example: $100^*\pi$ will be printed as:

F12.4	314.1593
E12.4	0.3142E+03
E12.4E3	0.3142E+003
EN12.4	314.1593E+00
ES12.4	3.1416E+02

Edit descriptors

- COMPLEX: edited using pairs of f, e, en, og es descriptors.
- LOGICAL: lw-descriptor.
- CHARACTER: a-descriptor:
 - A: width determined by the input/output.
 - Aw: width

Example: string: STATEMENT (9 chars)

A9: STATEMENT, A10: STATEMENTb, A8: STATEMEN

Edit descriptors

- General: g-descriptor may be used for any intrinsic data type.
 - gw.d
 - gw.dee
- When used for reals it is identical to the e-descriptor.
- The g-descriptor automatically chooses between the e and f descriptor.
- Derived types are edited by the appropriate sequence of edit descriptors corresponding to the intrinsic types of the type.

Edit descriptors

- Minimal field width editing:
 - If the width is set to zero (i0, f0.3 etc.) the minimal width will be used.
- Control edit descriptors:
 - Embedded blanks may be controlled by the bn (blanks null) or bz (blanks zero) descriptor.
Example: bb1b4 with bn,i5 returns 14 and
bz,i5 returns 104.

Edit descriptors

- Control edit descriptors:
 - Leading signs may be controlled by the ss (sign suppress) and sp (sign plus)
Example: 1000 with sp,en9.2 returns: +1.00E+03
 - Scaling factor: kp-editor, where k is a default integer literal constant.
Example: 1000 with 3p,e9.2 returns: 100.0E+01
 - Colon editing: will terminate format control if there are no further items in an I/O list.

```
program main
integer :: n
integer, dimension(5) :: l

n = size(l)
do i=1,n
    l(i) = 3*i
enddo
n = 3
print ' (" 11 = ",i4, : , &
      " 12 = ",i4, : , &
      " 13 = ",i4, : , &
      " 14 = ",i4, : , &
      " 15 = ",i4)' , (l(i),i=1,n)
end program main
```

Unformatted I/O

- Formatted I/O has a large overhead for the conversion and the file size is significantly larger than the corresponding binary representation of the data.
- Unformatted I/O:

Example:

```
REAL(MK), DIMENSION(100,100) :: array
```

```
OPEN(10,FILE='binary.dat',FORM='UNFORMATTED')
```

```
WRITE(10) array
```

or

```
READ(10) array
```

Direct-access files

- The default file access in FORTRAN is SEQUENTIAL – that is data can only be inserted at the end of the file.
- Direct-access files allows to insert or read an arbitrary record.

Operations on external files

External files

- FORTRAN is connected to a file through the *UNIT* number. A unit number must not be connected to more than one file at once, and a file must not be connected to more than one unit at once.
- Notice also:
 - The set of allowed names for a file is processor dependent.
 - A file never contains both formatted and unformatted records. (FORTRAN 2003 allows `ACCESS='STREAM'`)

File positioning

- The backspace statement:
 - `BACKSPACE([unit=]u[,iostat=ios][,err=error-label])`
Position the file before the current record.
This statement is often very costly in computer resources and should be used as little as possible.
- The REWIND statement
 - `REWIND([unit=]u[,iostat=ios][,err=error-label])`
Will rewind the file to the beginning of the file.

File positioning

- The ENDFILE statement
 - ENDFILE([unit=]u[,iostat=ios][,err=error-label])

Will write an End Of File (EOF) to the unit (thus effectively truncate the file).
- Data transfer statements:
 - READ
 - WRITE or PRINT

The OPEN statement

- The OPEN statement is used to connect an external file to a unit, create a file that is pre-connected, create a file and connect it to a unit, or change certain properties of a connection:

OPEN([unit=]u[,olist])

- iostat=ios, ios is zero on success.
- err=err-label: label of the statement in the same scoping unit to which control is transferred if an error occurs.

The OPEN statement

- status=st, where st is a character expression that provides the value: old, new, replace, scratch, or unknown. With “scratch” the file= is illegal and the file is deleted by the matching CLOSE([unit=]u).
- access=SEQUENTIAL or DIRECT.
- form=UNFORMATTED or FORMATTED (default).
- recl=rl the record length.
- blank=NULL or ZERO.
- position=ASIS, REWIND, or APPEND.
- action=READ, WRITE or READWRITE.
- delim=QUOTE, APOSTROPHE, or NONE (default).
- pad=YES (default) or NO.

The OPEN statement

Example:

```
OPEN(10,FILE='bin.dat',FORM='UNFORMATTED')
```

```
WRITE(10) array
```

```
OPEN(11,FILE='ascii.dat')
```

```
WRITE(11,*) array
```

The CLOSE statement

- The CLOSE statement is used to disconnect an external file to a unit:

`CLOSE([UNIT=]u[,iostat=ios][,err=err-label][,status=st])`

– Status='delete' or 'keep' (default)

Example:

How to remove an external file in FORTRAN:

`OPEN(10,FILE='ABORT')`

`CLOSE(10,STATUS='DELETE')`

The INQUIRE statement

- The INQUIRE statement is used to inquire about the status and attributes of a file.
 - INQUIRE([unit=]u,ilist) or
 - INQUIRE(file=fln, ilist)

Where u refers to the unit number and fln to the filename. The ilist options include:

- iostat
- exist=.TRUE. If the file exist – otherwise .FALSE.

The INQUIRE statement

- opened=.TRUE. If the unit or filename are opened.
- number=the unit number assigned to the file. It returns -1 if it is not assigned.
- named=true if the file has a name.
- name=nam the name of the file.
- access=SEQUENTIAL/DIRECT (char).
- sequential=YES or NO
- form=FORMATTED or UNFORMATTED
- formatted=YES or NO (char).

The INQUIRE statement

- recl=record length
- nextrec=the number of the last record read or written.
- blank=NULL or ZERO.
- position=REWIND, APPEND or ASIS.
- action=READ, WRITE or READWRITE.
- read=YES, NO or UNKNOWN.
- write=YES, NO or UNKNOWN.
- readwrite=YES, NO or UNKNOWN.
- delim=QUOTE, APOSTROPHE or NONE.
- pad=YES or NO.

Recommendations

Use STANDARD tools

- Use STANDARD FORTRAN 90/95 (move to FORTRAN 2003 if/when needed).
- Use STANDARD MPI/OpenMP.
- Use STANDARD cpp for source preprocessing.
- Use STANDARD make for compilation.
- USE STANDARD editors (vi, emacs, nano, gedit).
- Use STANDARD revision control systems (rcs,cvs/svn,git).

Compilation

- Read the *man* pages for the compiler.
- Initially compile with `-C` to check for array bound violation.
- When `-C` passes switch to `-O3` (or more).
- Profile the code using compiler instructions (`f90 -p ...` or similar) or through call to `CPU_TIME()`.
- Compile on different platforms and with different compilers during the development of the code.

Style

- Write all FORTRAN keywords in capital letters.
- Indent 3 spaces after branches, loops etc..
- Do NOT use the tab – use the spacebar !
- Limit the length of lines to maximum 80 characters.
- A name should not exceed 31 characters.
- Keep variables names short, but meaningful.
- Variable names of length one should ONLY be used for loop counters (fx.: i,j,k)
- Add comments in the source code to describe variables.
- RCS/CVS/SVN: consider using \$Log\$ in the header.

Coding

- Use IMPLICIT NONE
- Define the precision in an include file 'param.h' or module

```
INTEGER, PARAMETER :: MK = KIND(1.0E0)
```

- Declare the floats using the MK:

```
INCLUDE 'param.h'
```

```
REAL(MK) :: value
```

```
REAL(MK), DIMENSION(:), POINTER :: array
```

```
value = 1.0_MK
```

Coding – Implicit interface

- Use ONE source file for ONE subroutine.
- Encapsulate each subroutine in a MODULE.
- Will automatically build an interface – so you NEVER need to type things twice = you NEVER need to type INTERFACE blocks.
- Subroutines in modules can only be accessed by USE – so you are forced to use the (implicit) interface; with explicit interfaces you might forget.
- Declare ALL arrays local to a subroutine in the MODULE which encapsulates the subroutine.
- Arrays in MODULEs will be allocated on the HEAP – automatic/local arrays can/will be allocated on the STACK !

Coding – explicit interface

- Use ONE source file for ONE subroutine.
- Write an interface file for each routine.
- Use separate module files for each subroutine to hold local data or write the module in the FORTRAN file but without including the subroutine in the module (no CONTAINS)

```
MODULE m_sub
```

```
    REAL, DIMENSION(:), POINTER :: work
```

```
END MODULE m_sub
```

```
SUBROUTINE sub(x)
```

```
USE m_sub
```

```
REAL, DIMENSION(:), POINTER :: x
```

Coding

- Do NOT use AUTOMATIC arrays!
- Do NOT allocate LOCAL arrays – but put them in a MODULE.
- Use the ONLY and PRIVATE module options.
- ALWAYS check the status of ALLOCATE !
- Remember to DEALLOCATE.

Coding

- Use cpp to encapsulate hardware/library specific source:

```
#ifdef __LAPACK
```

```
    CALL LAPACK_ROUTINE(...)
```

```
#else
```

```
    CALL MY_BEST_VERSION(...)
```

```
#endif
```

- All cpp directives should start with # in column 1.
- Do not break FORTRAN character strings (this is usually incompatible with cpp).

Coding

- Never use quotes (‘ or “) unless you are defining a FORTRAN string (also not in comments); quotes will disturb cpp.
- Perform check of routine arguments.
- Define and check a return status from all internal and external routines.
- All routines should have “info” as the last non-OPTIONAL argument
- Every routine should have a jump-out label 9999 CONTINUE at its end.

Efficiency

- Do NOT use small arrays (vectors) – use registers:

- **BAD** performance:

```
com(1:2) = 0.0
```

```
DO i=1,N
```

```
    com(:) = com(:) + xp(:,i)
```

```
ENDDO
```

- **GOOD** performance:

```
comx = 0.0; comy = 0.0
```

```
DO i=1,N
```

```
    comx = comx + xp(1,i)
```

```
    comy = comy + xp(2,i)
```

```
ENDDO
```

Efficiency

- Use DO loops (do NOT use array notation).
- **BAD** performance:

$A = A + B$! Where A and B are arrays

- **GOOD** performance:

```
DO j=1,N
```

```
    DO i=1,M
```

```
        A(i,j) = A(i,j) + B(i,j)
```

```
    ENDDO
```

```
ENDDO
```

OR:

```
A(1:M,1:N) = A(1:M,1:N) + B(1:M,1:N)
```

Efficiency

- For maximum efficiency use UNFORMATTED I/O.
- Write the entire array (do NOT use DO loops or other array subsections)

Example of fast IO:

```
REAL, DIMENSION(10,11,12,13) :: array
```

```
OPEN(10,FILE='data',FORM='UNFORMATTED')
```

```
WRITE(10) KIND(array)
```

```
WRITE(10) SHAPE(array)
```

```
WRITE(10) array
```

```
CLOSE(10)
```

Extras

- Use `iargc()` and `getarg()` to read the command line arguments. This is NOT standard but virtually it is.

Example:

```
nargc= iargc()
```

```
DO i=0,nargc
```

```
CALL getarg(i,string)
```

```
PRINT*, 'argument: ',string
```

```
ENDDO
```

returns for `./a.out -a -b 23 input.file`

`./a.out`

`-a`

`-b`

`23`

`input.file`

Extras

- In Fortran 2003/2008 use the `COMMAND_ARGUMENT_COUNT()` and the `GET_COMMAND_ARGUMENT(number[,value][,length][,status])`

Example:

```
program main
character(len=256) :: arg
integer :: i,ilen,info,narg

narg = command_argument_count()
print*, 'command_argument_count() = ', narg
do i=1,narg
    call get_command_argument(i,arg,ilen,info)
    print*, 'arg = ', arg(1:ilen)
enddo
end program main
```

Submitting jobs

In HPC, simulations are usually submitting as a job to a queue:
The PBS queuing system format: qsub submit.sh:

```
#!/bin/sh
#PBS -N my_job_name
#PBS -l walltime=00:02:00
#PBS -l nodes=1:ppn=16
#PBS -m bea

# change directory to the directory from where you submitted the job from
cd $PBS_O_WORKDIR

# or change explicitly to a specific directory
cd /home/jhwa/runs/case001/

# run the executable:
./a.out > stdout
```


Submitting jobs

The SLURM queuing system: sbatch submit.sh:

```
#!/bin/bash
#SBATCH --job-name=case022
#SBATCH --no-requeue
#SBATCH --mail-type=ALL
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=24
#SBATCH --time=40:00:00
#SBATCH --output=mpi_job_slurm.log
#SBATCH --partition=xeon24
#-----
# change to local directory
#-----
cd $SLURM_SUBMIT_DIR

# or change explicitly to a specific directory
cd /home/jhwa/runs/case001/

# run the executable:
./a.out > stdout
```